


Theodore Wu, Charlie Liu, Alexa Daigle, Espie Taylor
Assignment 4
Parallel Programming
03/29/24

Total execution time for each run:

- A.) 1 node, 1 GPU, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.643961
- B.) 1 node, 2 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.656310
- C.) 1 node, 3 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.655127
- D.) 1 node, 4 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.646586
- E.) 1 node, 5 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.649848
- F.) 1 node, 6 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.659207
- G.) 2 nodes, 12 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5: 0.722893



Maximum speedup relative to using a single GPU:

Time with single GPU:

Speedup with 2: $0.643961/0.656310 = 0.981x$

Speedup with 3: $0.643961/0.655127 = 0.982x$

Speedup with 4: $0.643961/0.646586 = 0.996x$

Speedup with 5: $0.643961/0.649848 = 0.991x$

Speedup with 6: $0.643961/0.659207 = 0.977x$

Speedup with 12: $0.643961/0.722893 = 0.891x$

Maximum speedup: 0.996x

Configuration that yields the fastest "cells updates per second" rate:

A.) 1 node, 1 GPU: World of 16384^2 that runs for 128 iterations with 0.643961 execution time:

$16384^2 * 128 = 34,359,738,368$ cell updates

$34,359,738,368 / 0.643961 = 53,356,862,244.8$ cell updates per second

B.) 1 node, 2 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.656310 execution time:

$2(16384^2) * 128 = 68,719,476,736$ cell updates

$68,719,476,736 / 0.656310 = 104,3705,820,018$ cell updates per second

C.) 1 node, 3 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.655127 execution time:

$3(16384^2) * 128 = 103,079,215,104$ cell updates

$103,079,215,104 / 0.655127 = 157,342,339,888.3$ cell updates per second

D.) 1 node, 4 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.646586 execution time:

$4(16384^2) * 128 = 137,438,953,472$ cell updates

$137,438,953,472 / 0.646586 = 212,560,979,470.6$ cell updates per second

E.) 1 node, 5 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.649848 execution time:

$5(16384^2) * 128 = 171,798,691,840$ cell updates

$171,798,691,840 / 0.649848 = 264,367,501,077.2$ cell updates per second

F.) 1 node, 6 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.659207 execution time:

$6(16384^2) * 128 = 206,158,430,208$ cell updates

$206,158,430,208 / 0.659207 = 312,737,016,154.3$ cell updates per second

G.) 2 nodes, 12 GPUs/MPI ranks: World of 16384^2 that runs for 128 iterations with 0.722893 execution time:

$12(16384^2) * 128 = 412,316,860,416$ cell updates

$412,316,860,416 / 0.722893 = 570,370,525,673.9$ cell updates per second

Fastest "cells updates per second" rate:

2 nodes, 12 GPUs/MPI ranks, 16Kx16K world size each MPI rank, 128 iterations with 256 CUDA thread block size and pattern 5 with 570,370,525,673.9 cell updates per second

Why this configuration was faster than others:

Contributions:

All group members engaged in a collaborative effort to review and test the C and CUDA code to ensure correctness. Additionally, all group members effectively communicated with one another in regard to dividing up work and completing various tasks in a timely manner.

Other individual contributions:

Theodore Wu: Corrections and bug fixes in code

Charlie Liu: C and CUDA code

Alexa Daigle: Report and calculations

Espie Taylor: AIMOS runs